



Principled Futures

BOARD RESEARCH PAPER

The Boardroom Compass: Real-Time Metrics for AI Governance

Principled Futures Advisory — A Salveus Labs Product · July 2026

© 2026 Salveus Labs. All rights reserved.

principled futures · research

A real-time risk dashboard for AI governance represents a shift from static reporting to continuous telemetry, serving as a visual hub that translates technical data into actionable risk intelligence. This allows boards to oversee the "orchestration level" of autonomous systems rather than just individual tasks. Based on the research, specific indicators for a real-time dashboard are categorized as follows:

— 1. Visual Risk Indicators and Alerts

- Risk Heat Maps: Interactive visualizations that surface AI risks alongside traditional operational, financial, and compliance risks.
- Predictive Indicators: Indicators designed to detect emerging threats—such as security breaches or performance degradation—before they escalate into "balance-sheet events".
- Threshold Alerts: Predefined triggers for metrics like error rates or bias ratios that provide early warning alerts for immediate board awareness.

— 2. Model Health and Performance Telemetry

- Model Drift Monitoring: Real-time tracking of accuracy or relevance degradation as data patterns evolve.
- Anomaly Detection: Monitoring for unexpected or unexplainable outputs, which may indicate model failure or data poisoning.
- Reliability Metrics: Tracks system uptime, latency, and the frequency of data validation errors during model updates.
- Retraining Cadence: Time elapsed between model retraining cycles to ensure relevance.

— 3. Ethical and Fairness Telemetry

- Bias Indicators: Quantitative metrics such as demographic parity or disparate impact ratios to detect discriminatory outcomes in hiring, lending, or other high-stakes areas.
- Explainability Scores: The percentage of production AI models with documented explainability reports and the volume/response time of user requests for decision explanations.
- Inaccuracy and Hallucination Rates: Tracking the frequency of inaccurate or hallucinated outputs across generative AI systems.

●

— 4. Security and Resilience Telemetry

- Active Threat Monitoring: Real-time tracking of security vulnerabilities and results from adversarial attack simulations (e.g., prompt injection or model inversion attempts).
- Kill Switch Readiness: The status of "kill switch" protocols, showing whether dual-authorization mechanisms (typically involving the CISO and CTO) are ready to halt operations if unintended harm is detected.
- Shadow AI Detection: Scans the network for unauthorized AI deployments and unapproved public tools being used by employees.

— 5. Strategic and Value Realization Metrics

- ROI Tracking: Documented business value such as time saved, costs avoided, or revenue attributed to AI initiatives compared to original business cases.
- Strategic Alignment: Distribution of AI investments across efficiency-focused versus growth-focused use cases.
- Scaling Status: The percentage of AI use cases successfully moved from pilot programs to scaled operations in core functions.

— 6. Workforce and Cultural Indicators

- AI Fluency: Tracking the percentage of staff who have completed required responsible AI training and literacy programs.

- Ethical "Speak-Up" Reporting: The frequency and resolution of ethics-related concerns or incidents raised by employees or stakeholders.

- Employee Sentiment: Regular survey data regarding workforce perceptions of AI's impact on job security, fairness, and overall engagement.